

Where's my Ferry?

Support Vector Machines for Modeling Ferry Tardiness

Kyle Wenholz, advised by Professor Brad Richards

April 5, 2013

Contents

1	Introduction	2
2	Preliminaries	2
2.1	ANN and SVM	3
2.2	Basics of SVM	4
2.3	Kernel mapping	5
2.4	LibSVM	6
3	Predicting tardiness through data	6
3.1	Washington State’s ferries	7
3.2	Weather and final SVM format	8
4	Results	9
4.1	The first pass	10
4.2	General results	10
4.3	The problem with weather	11
5	Closing thoughts	12
5.1	Improvements to the data analysis	12
5.2	Extensions	13
5.3	Conclusion	13

1 Introduction

Ferries (as in **Figure 1**) present an interesting opportunity to examine a complex traffic system filled with data and affecting the many lives. This paper will discuss a new approach to predicting the timeliness of ferries using a support vector machine. The research was part of an Honors Program senior thesis at University of Puget Sound, supervised by Professor Brad Richards. An overview of how support vector machines work and their complications in real world use will introduce our discussion of how this powerful construct presents a powerful and intuitive model for tackling the massive data associated with ferries. A functioning model is developed to examine the complexities of this approach, the practicality, and as a means of exploring some of the data. Additionally, the hypothesis that the addition of weather information improves accuracy of the model is explored.



Figure 1: A prototypical WSDOT ferry on its route.

We attempt to predict the tardiness of ferries, where tardy is defined as no less than three minutes later than the initial (at day’s beginning) scheduled arrival or departure time. This definition of tardiness was chosen to allow for a non-trivial number of tardy events (more than 10% in the data) and to remain practical: i.e. someone could use the restroom in three minutes time or otherwise make use of that information. The Washington State Department of Transportation’s ferry system in the Pacific Northwest served as the source for all data. This system is of particular interest for its sheer volume of passengers: 22 million per year [Was12]. The volume lends itself well to the application of data mining techniques. While many techniques exist for data mining and prediction in general, support vector machines are used in this project for a variety of reasons.

2 Preliminaries

We consider two approaches popular in terms of both available implementations and the literature to be highly applicable to this problem: artificial neural networks (ANNs) and support vector machines (SVMs). These two approaches are algorithms, or machines, for *supervised learning*: labeled (known) data is fed into an algorithm that learns to make better predictions through comparing its own predictions to the labeled data set. Classical examples of using ANN and SVM are spam detection in emails, determining images in which cancer is present, or even detecting gender from a facial image. ANN and SVM are often considered to be very similar in the problems they tackle, especially since each can be used to train a linear classifier: a function taking in some object and determining a class to which it belongs. Rather than considering tardiness as a spectrum of degrees, we approach tardiness as a binary feature: late or not. This conceptually simplifies our problem and makes a

larger set of ANNs and SVMs available for use . If artificial neural networks and support vector machines can both solve classification problems, however, which are we supposed to choose?

2.1 ANN and SVM

Byvatov et al examine the differences of the two approaches for classifying drugs [BFSS03]. While their model certainly isn't for a ferry system, their discussion suggests data sets with many features (an aspect of the ferry model we discuss later) are better suited for SVM, and they also find little difference in the predictions of ANN and SVM. Most importantly, their conclusion, and that of their references, is the two approaches are complementary. Each tool catches different cases, with comparable results in each approach, and SVM may even require less “tweaking”. They concede that SVMs are generally easier to conceptualize, lending a greater sense of confidence to the non-expert.

Support vector machines are more intuitively motivated. Consider mapping the features of a ferry trip (time of departure, weather, boat name, etc.) into a coordinate plane. Now, if we do this for all the points in our training set, we can also attach labels to the points (late or not, as in **Figure 2**). In two dimensions, we could try drawing a line through the late and not-late points to separate them. This line is simply a hyperplane in higher dimensions, and SVM has its theoretical underpinnings based on the idea of drawing this hyperplane as an optimal separator. Apart from the possible ease of use and understanding, there are also examples of SVMs being used in similar traffic prediction systems, providing a baseline comparison and the final reason for choosing SVMs over artificial neural networks.




Figure 2: A simple example of a separating line for tardiness. Each data point (or instance) has features X_1 and X_2 , corresponding to axes on the graph.

Smith et al use SVM to predict the timeliness of airline traffic in conjunction with weather patterns [SSD08]. Their goal is to predict the likelihood of requiring a “ground delay program”: air traffic control shuffling flight times to account for tardiness in the schedule. The model they use involves traffic flow management programs to estimate the effects of capacity on an air traffic system, and the SVM is trained on labeled data to produce a function predicting the need of a ground delay program and/or an actual delay. Their system was found to be 78% accurate in predicting need of a ground delay program and 83% accurate in predicting a delay. Air traffic is much like the ferry

system given its sensitivity to weather, attempt to adhere to a strict schedule, and heavy use. Due to these similarities, Smith's access to resources like the AMPL supercomputer, and his group's expert knowledge, we use 80% as the target accuracy for our ferry model.

As a final note regarding the usefulness of SVMs, they can be used for regression analysis in large data sets with many variables [CL11]. While our project does not utilize this aspect of SVMs, the possibility of performing such an analysis with the same tool was one more reason for choosing support vector machines.

2.2 Basics of SVM

A brief overview of the principles behind support vector machines and their use will help to explain basics of this project. As stated earlier, SVMs are used to find a linear classifier (some function) for a set of data. In this project, we seek a way to separate ferry trips that will be three minutes past their scheduled time (late) or less (on time). To find the linear classifier, an SVM uses a feature vector, F , to describe ferry trips. An example F may have entries for

$$F = \begin{bmatrix} EstimatedArrival \\ BoatName \\ Temperature \\ WindSpeed \end{bmatrix}.$$

This feature vector needs to be entirely numerical, so categorical variables like boat name must be encoded as numbers. It is possible to assign boat names to different values, but most often, it is best to expand the feature like a bit vector. If there were boat names *SS Minnow*, *Death Star*, and *Millennium Falcon*, then our new F would be

$$F = \begin{bmatrix} EstimatedArrival \\ SS_Minnow \\ Death_Star \\ Millennium_Falcon \\ Temperature \\ WindSpeed \end{bmatrix},$$

where the entries for boat names are a 1 if this trip has that name and 0 otherwise. For reasons better detailed in [CL11] this is considered best practice, but suffice it to say that this format is better scaled, and leads to better SVM performance (both in runtime and accuracy).

Along with feature vectors for every trip, an SVM uses a weight vector, w , as the backbone for the linear classifier. Thus, this w is what the SVM trains. The dot product of F and w is taken to determine which class a trip's feature vector corresponds to, with negative outputs corresponding to late and positive to on time. Moreover, while training on the labeled data the SVM adjusts w whenever it predicts incorrectly. The amount to adjust can vary based on the implementation, but over time w improves on its ability to predict the class of any instance in the training set.

An optimal solution to w would form a hyperplane separating the instances perfectly by class. In **Figure 3** is an example of how an SVM might train a weight vector for a linear classifier (represented graphically as a line) over a set of training data. The final vector, H_3 , is considered optimal because it maximizes the distance between itself and the nearest points of each class with no two points of a

different class on the same side of the line. In the example, it is possible to separate the data with a line, but in the real world, and our ferry problem, this is rarely the case.




Figure 3: As in **Figure 2**, we are trying to separate labeled points with a line. An SVM does this iteratively, meaning line H_1 would be the first pass, H_2 the result of refitting w , and H_3 the final classifier produced by the SVM.

2.3 Kernel mapping

In **Figure 4** is an example of data we can't separate. Most real world data is like this: inseparable by a hyperplane. The problem this creates is that we can no longer look for a linear classifier. A nonlinear classifier is required, a significantly more difficult problem. We might consider looking for something within reasonable bounds, since no optimal hyperplane exists. It is difficult to formulate how long an SVM should “look” for the hyperplane, however. Fortunately, the kernel trick does away with this problem through a transformation of the data [ABR64].

The kernel trick maps the instances of data into a higher dimensional space (often much higher), and the SVM works to find a separating plane in the new space. The example in **Figure 4** gives an idea of what the mapping would do to the data. The kernel function used in this project is a Gaussian radial basis function and was chosen for its popularity and since it is built into the LibSVM library [CL11] (discussed in the next section).

Using the kernel trick requires a small addition to the training process. Optimal parameters for the kernel function need to be found before fitting with the SVM. This is done by an exhaustive grid search:

1. The user supplies a space of parameters to try.
2. Within this space, a search grid is created.
3. Each grid square is “tried” by using parameters within to transform the data and then run the SVM on a subset of the training data.
4. The best grid is then turned into a grid and the process repeats until results stabilize.




Figure 4: Separating the data in the first part requires a mapping into an often much higher dimension. This mapping makes it possible for a learning algorithm like SVM to find at least some separating hyperplane.

So, training the kernel parameters requires training the SVM on very small sets of data many times to find the best parameters in a space provided by the user.

Implementing all of the functionality of an SVM and kernel function would require expert knowledge and a thorough amount of validation. There are also concerns regarding optimality and an ability to compare results. Fortunately, libraries to work with SVMs and kernel functions already exist.

2.4 LibSVM

LibSVM [CL11] from Chih-Chung Chang and Chih-Jen Lin provides the implementation of SVMs we used. This library has bindings to many languages, is reasonably straightforward to use from the command-line, is well tested, and has an acceptable amount of documentation. In fact, their introduction guide [?] is incredibly helpful both for understanding how SVMs work and using their implementation. It is also very encouraging that the LibSVM implementation is cited in thousands of papers. Using this implementation of an SVM allowed a great deal of time to be saved and provided a higher degree of certainty in the results.

3 Predicting tardiness through data

Support vector machines require two sets of labeled data to find a linear classifier: training and testing. The training data is fed into the SVM as a means of training the weight vector and finding a linear classifier, and the testing data is used after this process to check the accuracy of the SVM's product. It is crucial that these two sets be disjoint. Otherwise, the SVM is basically cheating by being tested on data it already learned from. There is no clear indicator for how much data is needed to successfully train an SVM, but too little means the SVM can't generalize and too much can cause overfitting. Overfitting is where an SVM is trained on too much data of similar qualities so that it fails to accurately classify data of different qualities. This roughly corresponds to finding patterns which only apply to a small subset of the real instances. So, it often comes down to the engineer to determine the appropriate amounts of data for the training and testing sets based on experience, experimentation, and availability.



Figure 5: The Washington State Department of Transportation ferry system. 20 terminals, making 10 routes, serve users from Point Defiance in Tacoma up to Victoria in British Columbia [oT12c].

3.1 Washington State’s ferries

The Washington State Department of Transportation (WSDOT) runs the ferry system in the Pacific Northwest (since 1951) over a stretch of more than 130 miles, from Victoria, British Columbia to Point Defiance in Tacoma [oT12b]. The system serves over 22 million passengers per year in about 160,000 sailings [Was12]. We will refer to one sailing (from one terminal to another) as a trip. **Figure 5** gives the map of the WSDOT’s 20 terminals and 10 routes, lending perspective to the size of the entire system.

One concern during this process is that the ferry system itself is on time 95% of the time. It’s not clear how this impacts the practicality of an SVM in this problem nor what reasonable upper and lower bounds on expected accuracy should be. Because of this, the benchmark of 90% from the airlines paper [SSD08] remains the target in this project.

The system is interesting for more than just its traffic. Several of the routes are rather windy looking, suggesting they may be more susceptible to weather and traffic delays. Some terminals, furthermore, are a hub for several routes. We might conjecture these terminals to be susceptible to pileups in traffic. These and other complexities of the system make it worthwhile in pursuing the robust model provided by an SVM. Fortunately, the complexities and heavy use of the system seem to have convinced the WSDOT to keep thorough records of the traffic.

To effectively use an SVM, we need enough data we can experiment with and it needs to be in a uniform, accurate state. Originally, a web scraper was written to grab information from the WSDOT VesselWatch [oT12c] page, but this was made obsolete when a request for information was eventually fulfilled. The WSDOT provided, upon request, a record of approximately 340,000 ferry trips from October of 2010 to October of 2012. The data came as a comma separated value file with the nine fields. An example is shown in **Figure 6**.

The column headings are all fairly straightforward, but it is worth noting the scheduled and initial

Vessel	Departing	Arriving	Sched Depart	Actual Depart
Kittitas	Mukilteo	Clinton	9/1/2010 0:00	9/1/2010 0:01
Sealth	Vashon	Southworth	9/1/2010 0:05	9/1/2010 0:06
Tacoma	Colman	Bainbridge	9/1/2010 0:15	9/1/2010 0:19
Initial ETA	Actual Arrival	Date	Route Name	
9/1/2010 0:13	9/1/2010 0:14	9/1/2010	Mukilteo - Clinton	
9/1/2010 0:16	9/1/2010 0:19	9/1/2010	Southworth - Vashon	
9/1/2010 0:45	9/1/2010 0:51	9/1/2010	Seattle - Bainbridge Island	

Figure 6: The header of the data provided from the WSDOT (broken into two tables to fit page widths). There are over 340,000 lines in the original file.

values, for departure and arrival respectively, are the times found on the schedule for the beginning of the day. With the way the data is given, we separate out the times into departures and arrivals, creating two separate scheduling problems. This way it is possible to try predicting arrivals with the added information of departure time (a seemingly easier problem).

Before feeding the data into the SVM, it was necessary to transform times and categorical variables into numeric values. This was done through several Python scripts: using seconds past January 1, 1970 for the times and expanding the categorical out as many variables functioning like a bit vector. Additionally, and a far more difficult problem, we joined this ferry data with weather data obtained from the National Oceanic and Atmospheric Administration (NOAA).

3.2 Weather and final SVM format

Retrieving accurate and full historical weather data presented an unexpected challenge. Most weather sites are focused on future weather and make historical data available only through programming APIs which may require a paid subscription. Fortunately, NOAA provides (by subscription) a portal for retrieving weather data by station. Stations record over 20 different features at some point every hour. It seems that some stations experience outages or are unable to record on some instruments every so often; because of this, we used the station with the seemingly most complete data set in the area. This turned out to be a station near the Tacoma Narrows Bridge. The data was downloaded in files by month for the range of October 2010 to October 2012 from a quality controlled store on NOAA’s site [OA12]. Only about 1,000 of the more than 24,000 weather entries had to be removed for null values in the wind data. Unfortunately, the station is not a central point, but it seemed close enough at the time.

In **Figure 7** is an example of what the weather data looks like in its raw form. Note there are some categorical variables, but nearly everything is a numeric. The categorical variables weren’t utilized in the final model, since they can be inferred from the other data. Many of the columns are simply blank, due to the stations lack of equipment for certain measures. The measures for dry bulb, wet bulb, dew point, wind speed, wind direction, and station pressure were used in the final model. There was very little missing data for any of these, and none of them directly imply another. Humidity may seem to be missing, but it can be calculated by the dry bulb and wet bulb temperatures.

We avoided using multiple stations to keep the joining of ferry and weather data simpler: since all ferries matched with one station, we just found the weather recording closest (within at least an hour) of the actual time reported on the ferry trip. Of the over 340,000 ferry trips, only about

WBAN	Date	Time	StationType
SkyCondition	SkyConditionFlag	Visibility	VisibilityFlag
WeatherType	WeatherTypeFlag	DryBulbFahrenheit	DryBulbFahrenheitFlag
DryBulbCelsius	DryBulbCelsiusFlag	WetBulbFahrenheit	WetBulbFahrenheitFlag
WetBulbCelsius	WetBulbCelsiusFlag	DewPointFahrenheit	DewPointFahrenheitFlag
DewPointCelsius	DewPointCelsiusFlag	RelativeHumidity	RelativeHumidityFlag
WindSpeed	WindSpeedFlag	WindDirection	WindDirectionFlag
ValueForWindCharacter	ValueForWindCharacterFlag	StationPressure	StationPressureFlag
PressureTendency	PressureTendencyFlag	PressureChange	PressureChangeFlag
SeaLevelPressure	SeaLevelPressureFlag	RecordType	RecordTypeFlag
HourlyPrecip	HourlyPrecipFlag	Altimeter	AltimeterFlag

94274, 20110101, 0053, 12, *CLR*, , 10.00, , , 25, , -3.9, , 22, , -5.4, ,
16, , -8.9, , 69, , 0, , 000, , , 29.69, , 1, , 002, , 30.04, , *AA*, , ,
, 30.03,

Figure 7: An example line from a weather data file preceded by the many, but thankfully explicit, column headings.

13,000 had to be removed for lack of weather data. With a data set so large, this only constituted 4%, an acceptable loss in our case. To join the data, a simple Python script searched the weather data for a nearest time to join with each ferry trip, and then the departure and arrival information were separated to place each as an individual event in separate files. It was then straightforward to calculate the label (late or not) for each event by subtracting the appropriate times.

Following the joining of weather and ferry data and the labeling, it is necessary scale the data between -1 and 1 , since this is considered an SVM best practice to make no feature stand out by sheer magnitude [?]. LibSVM performs this transformation quickly and is also useful for separating out the training data from the testing data. We used it to randomly sample about 290,000 data points for training and 30,000 for testing. The sampling was random to the extent that both sets of data contained a roughly equal distribution of late and not-late ferry trips. This is an incredibly important feature to preserve, since we want to both learn about tardiness and check if the linear classifier produced is able to predict both scenarios. With separated data in hand, we proceeded to train a linear classifier using the LibSVM package.

4 Results

The results for the project exist in three parts. As described before, the primary goal is to model the ferry system and answer the question of whether or not using an SVM to train a linear classifier is useful. This leads to the first two phases which involve a first pass and then a serious refining of the methodology. As a benchmark, we use approximately 90% accurate on the test data to be comparable with Smith et al [SSD08], since like that group, we wish to see if using SVM is a practical method for determining ferry tardiness. Recall that a secondary goal is to find which features matter in determining tardiness. While this goal is asking questions best answered by a regression analysis, we will attempt a crude approximation through empirical methods. The final component is a small experiment to test interesting phenomena arising in the general results. This latter part is possibly the most interesting of the project, but we begin with the first attempt to use LibSVM.

4.1 The first pass

We began by determining whether or not a ferry is on time for its departure, leaving arrival for after the process was worked out. A file with all of the labeled training instances was fed into LibSVM, using the default parameters of the kernel. The process took an entire night running on low end hardware, but the results were disappointing to say the least.

After running on the test data, the linear classifier only yielded 60% accuracy. This is significantly beneath the benchmark. The time to run the process was also surprising and led to a more cautious approach for the rest of the work. To improve the performance of the classifier, LibSVM provides an exhaustive guide [?] for beginners in the perils of simply “running it”. The directions weren’t followed on the first pass because it was hoped the process would be much simpler just using the default parameters.

Having learned some lessons, the LibSVM guide was followed to the letter on a second pass. The most significant change to the process was training on the parameters for the kernel function. As mentioned before LibSVM provides code to perform an exhaustive search for the best fit parameters to use in the function. The subsets for this training come from the training data. Finding the best fit parameters is a long process, sometimes as long as the actual SVM training, so it slows down the process significantly, but as we’re about to discuss the dividends are well worth it.

4.2 General results

Training the SVM to predict arrival time tardiness proved to be the most successful. In this case, a full instance included features for the actual departure time, all weather, predicted arrival time, date, vessel name, departure location, and arrival location. (Final results for all runs are displayed in **Figure 8**.) It is expected for this to be the problem the SVM can be most accurate about, since really it only needs to understand how long the trip will likely take a ferry.

The tardiness of full feature instances were able to be predicted with $\approx 87.0\%$ accuracy. For a second pass, this is incredibly close to the benchmark of 90%. To tune the linear classifier from our SVM, it is necessary to try trimming unnecessary features to find maximum accuracy. Due to time constraints, the trimming we perform here is very coarse. A thorough examination would attempt to take examine all combinations of features, but this is beyond the scope of our work, especially in regards to available computing power.

Removing the departure information from the features drops accuracy to 60%, a disappointing result. It may be possible to dramatically improve this result with some of the techniques discussed in **Section 5.1**. To determine whether removing more features made any significant difference, we also removed the weather information from predicting arrivals. This is, by far, the more interesting result.

A full departure feature vector includes all of the information regarding a trip (vessel name, time, route) and the weather. The prediction results here were not as exciting as with arrivals, but they proved to be more stable. The full departure vectors were predicted with 76.3% accuracy, and removing weather only reduced accuracy to 75.7%. While 76% is not at the benchmark, having experienced the change removing features and adjusting the training sets can have, it seems like more time with the data could yield even better results.

When departure information is still included in arrival predictions, the removal of weather actually improves accuracy resulting in 88%, an improvement of one percentage point. Without the departure

Arrival Predictions	
Configuration	Percent accuracy on test data
Full	87%
No Departure	60%
No Weather	88%
No Departure, No Weather	59.4%

Figure 8: Feature configurations and their associated percent accuracies in predicting the timeliness of arrivals convey the empirical importance of the features included.

Departure Prediction Accuracy	
Configuration	Percent accuracy on test data
Full	76.3%
No Weather	75.7%

Figure 9: Feature configurations and their associated percent accuracies in predicting the timeliness of departures convey the empirical importance of the features included.

information, the percent accuracy decreases by only 0.6 percentage points, a negligible difference at 59.4% accuracy. Weather was included as all but a given for improving the accuracy of the SVM work, but the results just outlined call into question this assumption and warranted further investigation.

4.3 The problem with weather

Since weather made barely any difference in prediction accuracy (even degrading it in some cases), examining possible causes of this unexpected incident became the last focus of the project. By re-examining the assumptions made regarding weather at the start, there are several which we can now dispute:

1. Weather will help improve accuracy.
2. The area is geographically small enough that one weather station should be just fine for all routes.
3. The weather variables dry bulb, wet bulb, dew point, wind speed, wind direction, and station pressure are the right variables to examine.
4. The data retrieved is of a sufficient quality.

Given time constraints and the results of the SVM, it seems assumption 2 is the best to tackle. But what how can we tease apart this assumption? Our strategy was to tackle the effect of the station location by comparing the effect of weather on three routes sufficiently far apart by running individual SVMs for each route (testing accuracy only on that route's trips). The three routes chosen were Tahlequah-Point Defiance, Kingston-Edmonds, and Coupeville-Port Townsend. Looking at **Figure 5**, we can see the furthest route, Coupeville-Port Townsend, is almost 100 miles from the weather station near Point Defiance. These three were deemed to be closest, middlemost, and furthest from the weather station, respectively, but also have enough trips to warrant an individual SVM.

The results of the experiment are mixed and can be viewed in full with **Figure 10**. We used full feature vectors for predicting both arrivals and departures, only taking out weather where noted for the comparison. We can see that the removal of weather rarely made a difference in any of the measures, except for the Kingston-Edmonds route when looking at arrivals. In that case the removal

Prediction accuracies with a <i>Weather/NoWeather</i> format	
Arrivals	Departures
Tahlequah-Point Defiance	
95.2%/93.9%	76.1%/75.3%
Kingston-Edmonds	
85.2%/94.7%	82.7%/82.8%
Coupeville-Port Townsend	
88.3%/87.7%	68.8%/69.4%

Figure 10: The arrival and departure time prediction accuracies for individual routes when including and excluding weather.

of weather was a major boon the accuracy. Otherwise, the results here are inconclusive of the actual efficacy of weather in improving predictions, and if anything, only help confuse the situation.

One final possible explanation is the robustness of this ferry system. The WSDOT reports their system to be on time 95% of the time [oT12a], a remarkable feat. It may simply be that there isn't much of a pattern in regards to weather or other predictable phenomena. We may be looking at a system whose only fault is vessel maintenance, a feature requiring very different data than that which we had access to.

The accuracy of the linear classifiers produced by our SVM runs are promising. With only novice level knowledge and hardware, we have improved significantly on random guesses for tardiness, and more time might see even better improvements.

5 Closing thoughts

The original goal of the project, to develop a ferry system model using AI techniques, was met and spawned further investigation into the behavior of the system. We nearly met the benchmark of 90% accuracy for a linear classifier produced by an SVM, and we concluded the performance could probably be improved on given more time and experience tuning the process. The exploration of weather's relationship to the ferry system also proved to be more complex than expected, clouding the assumption of there being a significant link between the two. The remaining questions and further work are now even greater in scope than when the project began.

5.1 Improvements to the data analysis

Most of the percent accuracies we have analysed were short of the 90% benchmark, but some were very close. More time to tune the features used in prediction, check the data quality, and experiment with different training set sizes would likely yield even better predictions. The latter goal is of particular interest. Since no definitive guide to using SVM was found before finishing the project, the size of training sets was kept the same throughout but may be a poorly informed choice. A common criticism of SVMs, and AI in general, is that training on too large of a dataset can lead to overfitting without realising it. This means the machine is only really capable of predicting events similar to the training data and can't generalize beyond. With a chance to train on smaller sets of data, we might find overfitting hampers the accuracy we found with larger training sets.

Recall the variance found among training an SVM over individual routes in the weather investigation. We discovered that individual routes varied widely, suggesting a "master" SVM for all routes may not be appropriate. We would like to assume that an SVM could tease apart the relationships

between routes and tardiness, but there may be too many conflicts between routes in terms of which features matter. For instance, a short route between Clinton and Mukilteo covers mostly open water and might not be very susceptible to weather in comparison to the island crowded route from Anacortes to Lopez. A thorough comparison of how individual SVMs perform for the routes might make this approach even more appealing or lend insight into a better choice.

5.2 Extensions

The end product of our SVM trainings worked in the ways we hoped: seeming to be useful for predicting ferry timeliness and as an investigative tool for the WSDOT ferry system. Still, the project is far from complete.

The SVM training produced useful linear classifiers, but a framework for average users could still be built. Whether this would be a web or mobile application is unfortunately restricted by the use of LibSVM. Since it is only implemented in C, a web application querying some server seems like the most straightforward approach. There are other possibilities since LibSVM has bindings in many languages, but the production of a real user interface and seeing it used would certainly validate the usefulness of predicting ferry times.

The SVM training itself could also be restructured. Several times, the lack of production time inhibited the ability to fully explore the data, but beyond this we could improve the definition of tardy. A multiclass SVM allows us to have degrees of tardiness (e.g. three minutes, five minutes, or ten minutes late). This sort of SVM is supported by LibSVM and extends fairly naturally from the two class case. Moreover, LibSVM could be used for a regression analysis of the features. This might be the best way to determine which features matter in regards to tardiness.

5.3 Conclusion

The use of an SVM for predicting ferry times novelly makes use of the large data sets available. This tool from artificial intelligence enables us to predict, with over 75% accuracy in most cases, the tardiness of ferries in the complex system, and it brings us insight into how important various features are when predicting this tardiness. While we may not fully understand the strange behavior exhibited by including weather in the model, the base assumption of weather's importance no longer appears valid. Further tuning of the SVM process is likely to yield even greater insight and perhaps better prediction results.

References

- [ABR64] A Aizerman, Emmanuel M Braverman, and LI Rozoner, *Theoretical foundations of the potential function method in pattern recognition learning*, Automation and remote control **25** (1964), 821–837.
- [BFSS03] Evgeny Byvatov, Uli Fechner, Jens Sadowski, and Gisbert Schneider, *Comparison of support vector machine and artificial neural network systems for drug/nondrug classification*, Journal of Chemical Information and Computer Sciences **43** (2003), no. 6, 1882–1889.
- [CL11] Chih-Chung Chang and Chih-Jen Lin, *Libsvm: a library for support vector machines*, ACM Transactions on Intelligent Systems and Technology (TIST) **2** (2011), no. 3, 27.
- [OA12] National Oceanic and Atmospheric Administration, *NCDC: Quality Controlled Local Climatological Data*, <http://cdo.ncdc.noaa.gov/qc1cd/QCLCD>, October 2012.
- [oT12a] Washington State Department of Transportation, *A Comparison of Operational Performance: Washington State Ferries to Ferry Operators Worldwide*, <http://www.wsdot.wa.gov/Research/Reports/700/750.1.htm>, November 2012.
- [oT12b] ———, *Washington State Ferries Our Fleet*, <http://www.wsdot.wa.gov/Ferries/yourwsf/ourfleet/>, October 2012.
- [oT12c] ———, *WSDOT - Ferries - VesselWatch*, <http://www.wsdot.com/ferries/vesselwatch/Default.aspx>, January 2012.
- [SSD08] David A Smith, Lance Sherry, and G Donohue, *Decision support tool for predicting aircraft arrival rates, ground delay programs, and airport delays from weather forecasts*, International Conference on Research in Air Transportation (ICRAT-2008), 2008.
- [Was12] Washington State Department of Transportation, *WSF Traffic Statistics*, http://www.wsdot.wa.gov/ferries/traffic_stats/, October 2012.