

Where's my Ferry?

Support Vector Machines for Modeling Ferry Tardiness

Kyle Wenholz, advised by Professor Brad Richards

April 21, 2013

Contents

1	Introduction	2
2	Preliminaries	2
2.1	ANN and SVM	3
2.2	Basics of SVM	4
2.3	Kernel mapping	5
2.4	LibSVM	6
3	Predicting tardiness through data	6
3.1	Washington State's ferries	7
3.2	Weather and final SVM format	9
4	Results	10
4.1	The first pass	11
4.2	General results	11
4.3	The problem with weather	12
4.4	Beyond accuracy	13
5	Closing thoughts	16
5.1	Improvements to the data and model analysis	17
5.2	Extensions	18
5.3	Conclusion	18

1 Introduction

Ferries (as in **Figure 1**) present an interesting opportunity to examine a complex traffic system filled with data and affecting many lives. This paper discusses a new approach to predicting the timeliness of ferries using a support vector machine. The research was part of an Honors Program senior thesis at University of Puget Sound, supervised by Professor Brad Richards (many thanks). An overview of how support vector machines work and their complications in real world use introduces our discussion of how this powerful construct presents a powerful and intuitive model for tackling the massive data associated with ferries. We develop a functioning model to examine the complexities of this approach, the practicality, and as a means of exploring some of the data. Additionally, we explore the hypothesis that the addition of weather information improves accuracy of the model.



Figure 1: A prototypical WSDOT ferry on its route.

We attempt to predict the tardiness of ferries, defining tardy as no less than three minutes later than the initial (at day’s beginning) scheduled arrival or departure time. We consider both arrival and departure times since we predict them separately later on. This definition of tardiness was chosen to allow for a non-trivial number of tardy events (more than 10% in the data) and to remain practical: i.e. someone could use the restroom in three minutes time or otherwise make use of that information. The Washington State Department of Transportation’s ferry system in the Pacific Northwest served as the source for all data. This system is of particular interest for its sheer volume of passengers: 22 million per year [Was12]. The volume lends itself well to the application of data mining techniques. While many techniques exist for data mining and prediction in general, support vector machines are used in this project for a variety of reasons.

2 Preliminaries

We consider two approaches popular in terms of both available implementations and the literature to be highly applicable to this problem: artificial neural networks (ANNs) and support vector machines (SVMs). These two approaches are algorithms, or machines, for *supervised learning*: feed labeled (known) data into an algorithm that learns to make better predictions through comparing its own predictions to the labeled data set. Classical examples of using ANNs and SVMs are spam detection in emails, determining images with cancer present, or even detecting gender from a facial image. ANNs and SVMs are often considered to be very similar in the problems they tackle, especially since each can be used to train a linear classifier: a function taking in some object and determining a class to which it belongs. Specifically, a linear classifier uses a linear combination of the object’s features,

often represented as a vector and called a feature vector, to make a classification decision.

Rather than considering tardiness as a spectrum of degrees, we approach tardiness as a binary feature: late or not. This conceptually simplifies our problem and makes a larger set of ANNs and SVMs available for use and adds questions regarding which features are important in the ferry system. If artificial neural networks and support vector machines can both solve classification problems, however, which are we supposed to choose?

2.1 ANN and SVM

Byvatov et al. examine the differences of the two approaches for classifying drugs [BFSS03]. While their model certainly isn't for a ferry system, their discussion suggests data sets with many features (an aspect of the ferry model we discuss later) may be predicted with a smaller standard error by SVM, albeit only marginally in some cases. Most importantly, they conclude, along with their references, the two approaches are complementary. Each tool yields different false positives (late), false negatives (not-late), true positives, and true negatives, although SVM may require less "tweaking". Even with such strong similarities, a few reasons exist, leading to the choice of SVMs.

To the non-expert support vector machines are intuitively motivated. Consider mapping the features of a ferry trip (time of departure, weather, boat name, etc.) into a coordinate plane. Now, if we do this for all the points in our training set, we can also attach labels to the points (late or not, as in **Figure 2**). In two dimensions, we could try drawing a line through the late and not-late points to separate them. This line is simply a hyperplane in higher dimensions, and SVM has its theoretical underpinnings based on the idea of drawing this hyperplane as an optimal separator. Apart from the possible ease of use and understanding, there are also examples of SVMs being used in similar traffic prediction systems, providing a baseline comparison and the final reason for choosing SVMs over artificial neural networks.




Figure 2: A simple example of a separating line for tardiness. Each data point (or instance) has features X_1 and X_2 , corresponding to axes on the graph.

Smith et al. use SVM to predict the timeliness of airline traffic in conjunction with weather patterns [SSD08]. Their goal is to predict the likelihood of requiring a "ground delay program": air traffic control shuffling flight times to account for tardiness in the schedule. The model they

use involves traffic flow management programs to estimate the effects of capacity on an air traffic system, and train the SVM on labeled data to produce a function predicting the need of a ground delay program and/or an actual delay. Their system was found to be 78% accurate in predicting need of a ground delay program and 83% accurate in predicting a delay. Combined with the discussion in [BFSS03], it seems that linear classifiers near 80% accuracy are within a range of acceptance.

Air traffic and the ferry system are quite alike: exhibiting a sensitivity to weather, attempt to adhere to a strict schedule, and heavy use. Due to these similarities, Smith’s access to resources like the AMPL supercomputer, and his group’s expert knowledge, we use 80% as the target accuracy for our ferry model. Accuracy serves as a crude investigative tool and measure of comparison for our first passes, but in **Section 4.4** we investigate the precision and recall of the various linear classifiers our work yields. These measures weren’t discussed in [SSD08] but grant us deeper insight into the usefulness of our work.

As a final note regarding the usefulness of SVMs, they can be used for regression analysis in large data sets with many variables [CL11]. While our project does not utilize this aspect of SVMs, the possibility of performing such an analysis with the same tool was one more reason for choosing support vector machines.

2.2 Basics of SVM

A brief overview of the principles behind support vector machines and their use helps to explain basics of this project. As stated earlier, SVMs are used to find a linear classifier (some function) for a set of data. In this project, we seek a way to separate ferry trips that will be three minutes past their scheduled time (late) or less (on time). To find the linear classifier, an SVM uses a feature vector, F , to describe ferry trips. An example F may have entries for

$$F = \begin{bmatrix} EstimatedArrival \\ BoatName \\ Temperature \\ WindSpeed \end{bmatrix}.$$

This feature vector needs to be entirely numerical, so categorical variables like boat name must be encoded as numbers. It is possible to assign boat names to different values, but most often, the best method expands the feature like a bit vector. If there were boat names *SS Minnow*, *Death Star*, and *Millennium Falcon*, then our new F would be

$$F = \begin{bmatrix} EstimatedArrival \\ SS_Minnow \\ Death_Star \\ Millennium_Falcon \\ Temperature \\ WindSpeed \end{bmatrix},$$

where the entries for boat names are a 1 if this trip has that name and 0 otherwise. Reasons better detailed in [CL11] consider this a best practice, but suffice it to say that this format is better scaled, and leads to superior SVM performance (both in runtime and accuracy).

Along with feature vectors for every trip, an SVM uses a weight vector, w , as the backbone for

the linear classifier. Thus, this w is what the SVM trains. Taking the dot product of F and w determines which class a trip's feature vector corresponds to, with negative outputs corresponding to late and positive to on time. Moreover, while training on the labeled data the SVM adjusts w whenever it predicts incorrectly. The amount to adjust can vary based on the implementation, but over time w improves on its ability to predict the class of any instance in the training set.

An optimal solution to w would form a hyperplane separating the instances perfectly by class. **Figure 3** illustrates an example of how an SVM might train a weight vector for a linear classifier (represented graphically as a line) over a set of training data. The final vector, H_3 , is optimal because it maximizes the distance between itself and the nearest points of each class with no two points of a different class on the same side of the line. In the example, we can separate the data with a line, but in the real world, and our ferry problem, this rarely works.




Figure 3: As in **Figure 2**, we are trying to separate labeled points with a line. An SVM does this iteratively, meaning line H_1 would be the first pass, H_2 the result of refitting w , and H_3 the final classifier produced by the SVM.

2.3 Kernel mapping

Figure 4 has an example of data we can't separate. This exemplifies most real world data: inseparable by a hyperplane. The problem this creates is that we can no longer look for a linear classifier. A nonlinear classifier must be used, a significantly more difficult problem. We might consider looking for something within reasonable bounds, since no optimal hyperplane exists. Difficulty arises in formulating how long an SVM should “look” for the hyperplane, however. Fortunately, the kernel trick does away with this problem through a transformation of the data [ABR64].

The kernel trick maps the instances of data into a higher dimensional space (often much higher), and the SVM works to find a separating plane in the new space. The example in **Figure 4** gives an idea of what the mapping would do to the data. The kernel function used in this project is a Gaussian radial basis function and was chosen for its popularity and since the LibSVM library implements it as the default kernel [CL11] (discussed in the next section).

Using the kernel trick requires a small addition to the training process. Optimal parameters for the kernel function need to be found before fitting with the SVM. An exhaustive grid search performs




Figure 4: Separating the data in the first part requires a mapping into an often much higher dimension. This mapping makes it possible for a learning algorithm like SVM to find at least some separating hyperplane.

this task:

1. The user supplies a space of parameters to try.
2. Within this space, create a search grid.
3. “Try” each grid square by using parameters within to transform the data and then run the SVM on a subset of the training data.
4. The best grid square turns into a grid and the process repeats until results stabilize.

So, training the kernel parameters requires training the SVM on very small sets of data many times to find the best parameters in a space provided by the user.

Implementing all of the functionality of an SVM and kernel function would require expert knowledge and a thorough amount of validation. There are also concerns regarding optimality and an ability to compare results. Fortunately, libraries to work with SVMs and kernel functions already exist.

2.4 LibSVM

LibSVM [CL11] from Chih-Chung Chang and Chih-Jen Lin provides the implementation of SVMs we used. This library has bindings to many languages, is reasonably straightforward to use from the command-line, is well tested, and has an acceptable amount of documentation. In fact, their introduction guide [CL11] is incredibly helpful both for understanding how SVMs work and using their implementation. It is also very encouraging that many thousands of papers cite the LibSVM implementation. Using this implementation of an SVM allowed a great deal of time to be saved and provided a higher degree of certainty in the results.

3 Predicting tardiness through data

Before we dig into the details of our data set, let us note a few details regarding the use of data with an SVM. Support vector machines require two sets of labeled data to find a linear classifier: training and testing. Feeding the training data into an SVM helps to adjust the weight vector for the linear

classifier, and the testing data helps to check the accuracy of the SVM’s product. These two sets absolutely must be disjoint. Otherwise, the SVM basically cheats by being tested on data it already learned from. There is no clear indicator for how much data is needed to successfully train an SVM, but too little means the SVM can’t generalize and too much can lead to *overfitting*. Overfitting is when an SVM trains on too much data of similar qualities so that it fails to accurately classify data of different qualities. This roughly corresponds to finding patterns which only apply to a small sample of the real instances, so it often rests on the engineer to determine the appropriate amounts of data for training and testing sets based on experience, experimentation, and availability. Below, we discuss our decision for these sets, primarily based on having practical seeming numbers of late and not-late instances in both sets along with the intent to experiment with various sizes later.

We are now ready to discuss the data used in the analysis and model creation: just how much is necessary, where it comes from, the formats, and the ways we processed it for the SVM training.

3.1 Washington State’s ferries

The Washington State Department of Transportation (WSDOT) runs the ferry system in the Pacific Northwest (since 1951) over a stretch of more than 130 miles, from Victoria, British Columbia to Point Defiance in Tacoma [oT12a]. The system serves over 22 million passengers per year in about 160,000 sailings [Was12]. We refer to one sailing (from one terminal to another) as a trip. **Figure 5** gives the map of the WSDOT’s 20 terminals and 10 routes, lending perspective to the size of the entire system.

One concern during this process is that the ferry system itself is on time 95% of the time. It’s not clear how this impacts the practicality of an SVM in this problem nor what reasonable upper and lower bounds on expected accuracy should be, so the benchmark of 80% from the airlines paper [SSD08] remains the target in this project. The WSDOT’s definition of on-time isn’t made clear in their report. Examining the data, however, reveals that our definition of on-time leads to approximately 13% of trips being late. For better or worse, this predictability in the system suggests a basic model could call everything not-late and be 87% accurate, but the discussion in **Section 4.4** shows that recall and precision are other important factors we might be able to improve on.

The system stands out for more than just its traffic. Several of the routes are rather windy looking, suggesting they may be more susceptible to weather and traffic delays. Some terminals, furthermore, are a hub for several routes. We might conjecture these terminals to be susceptible to pileups in traffic. These and other complexities of the system make the use of a complex SVM model attractive, and perhaps preferable given the vast data the WSDOT has stockpiled for their system.

To effectively use an SVM, we need enough data we can experiment with and it needs to be in a uniform, accurate state. Originally, a web scraper was written to grab information from the WSDOT VesselWatch [oT12b] page, but this was rendered obsolete when a request for information was fulfilled. The WSDOT provided, upon request, a record of approximately 340,000 ferry trips from September of 2010 to September of 2012. The data came as a comma separated value file with the nine fields. See **Figure 6** for an example.

The column headings are all fairly straightforward, but the scheduled and initial values, for departure and arrival respectively, are the times found on the schedule for the beginning of the day. We didn’t expect to have both predicted and actual times for both arrivals and departures, so the problem was reformulated to apply to tardiness for either type of scheduled time. To account for




Figure 5: The Washington State Department of Transportation ferry system. 20 terminals, making 10 routes, serve users from Point Defiance in Tacoma up to Victoria in British Columbia [oT12b].

Vessel	Departing	Arriving	Sched Depart	Actual Depart
Kittitas	Mukilteo	Clinton	9/1/2010 0:00	9/1/2010 0:01
Sealth	Vashon	Southworth	9/1/2010 0:05	9/1/2010 0:06
Tacoma	Colman	Bainbridge	9/1/2010 0:15	9/1/2010 0:19
Initial ETA	Actual Arrival	Date	Route Name	
9/1/2010 0:13	9/1/2010 0:14	9/1/2010	Mukilteo - Clinton	
9/1/2010 0:16	9/1/2010 0:19	9/1/2010	Southworth - Vashon	
9/1/2010 0:45	9/1/2010 0:51	9/1/2010	Seattle - Bainbridge Island	

Figure 6: The header of the data provided from the WSDOT along with the first three lines (broken into two tables to fit page widths). There are over 340,000 lines in the original file.

this, we separate the table in two: one for data pertaining to arrivals and one for departures. The arrival table actually includes the departure information since this is still pertinent to predicting the arrival time (an apparently easier problem with this information, which we discuss in **Section 4.2**). The two tables essentially give us two data sets and two problems: predicting arrival tardiness and predicting departure tardiness.

Before feeding the data into the SVM, it was necessary to transform times and categorical variables into numeric values. This was done through several Python scripts: using seconds past January 1, 1970 for the times and expanding the categoricals out as many variables functioning like a bit vector (addressed in **Section 2.2**). Additionally, and a far more difficult problem, we joined this ferry data with weather data obtained from the National Oceanic and Atmospheric Administration (NOAA).

3.2 Weather and final SVM format

Retrieving accurate and full historical weather data presented an unexpected challenge. Most weather sites are focused on future weather and make historical data available only through programming APIs which may require a paid subscription. Fortunately, NOAA provides (by a subscription our institution already has) a portal for retrieving weather data by station. Stations record over 20 different features at some point every hour. It seems that some stations experience outages or are unable to record on some instruments every so often; because of this, we used the station with the seemingly most complete data set in the area. This turned out to be a station near the Tacoma Narrows Bridge. The data was downloaded in files by month for the range of September 2010 to September 2012 from a quality controlled store on NOAA’s site [OA12]. Only about 1,000 of the more than 24,000 weather entries had to be removed for null entries. Unfortunately, the station is not a central point, but it seemed close enough at the time.

In **Figure 7** is an example of the weather data in its raw form. Note there are some categorical variables, but nearly everything is a numeric. The categorical variables weren’t utilized in the final model, since they can be inferred from the other data. Many of the columns are simply blank, due to the stations lack of equipment for certain measures. The measures for dry bulb, wet bulb, dew point, wind speed, wind direction, and station pressure were used in the final model. There was very little missing data for any of these, and none of them directly imply another. Humidity may seem to be missing, but it can be calculated by the dry bulb and wet bulb temperatures.

We avoided using multiple stations to keep the joining of ferry and weather data simpler: since all ferries matched with one station, we just found the weather recording closest (within at least an hour) of the actual time reported on the ferry trip. Of the over 340,000 ferry trips, only about 13,000 had to be removed for lack of weather data. With a data set so large, this only constituted 4%, an acceptable loss in our case. To join the data, a simple Python script searched the weather data for a nearest time to join with each ferry trip, and then the departure and arrival information were separated to place each as an individual event in separate files. It was then straightforward to calculate the label (late or not) for each event by subtracting the appropriate times.

Following the joining of weather and ferry data and the labeling, it is necessary to scale the data between -1 and 1 , an SVM best practice to make no feature stand out by sheer magnitude [CL11]. LibSVM performs this transformation quickly and is also useful for separating out the training data from the testing data. We used it to randomly sample about 270,000 data points for training and 50,000 for testing. The sampling was random to the extent that both sets of data contained a roughly

WBAN	Date	Time	StationType
SkyCondition	SkyConditionFlag	Visibility	VisibilityFlag
WeatherType	WeatherTypeFlag	DryBulbFahrenheit	DryBulbFahrenheitFlag
DryBulbCelsius	DryBulbCelsiusFlag	WetBulbFahrenheit	WetBulbFahrenheitFlag
WetBulbCelsius	WetBulbCelsiusFlag	DewPointFahrenheit	DewPointFahrenheitFlag
DewPointCelsius	DewPointCelsiusFlag	RelativeHumidity	RelativeHumidityFlag
WindSpeed	WindSpeedFlag	WindDirection	WindDirectionFlag
ValueForWindCharacter	ValueForWindCharacterFlag	StationPressure	StationPressureFlag
PressureTendency	PressureTendencyFlag	PressureChange	PressureChangeFlag
SeaLevelPressure	SeaLevelPressureFlag	RecordType	RecordTypeFlag
HourlyPrecip	HourlyPrecipFlag	Altimeter	AltimeterFlag

94274, 20110101, 0053, 12, *CLR*, , 10.00, , , 25, , -3.9, , 22, , -5.4, ,
16, , -8.9, , 69, , 0, , 000, , , 29.69, , 1, , 002, , 30.04, , *AA*, , ,
, 30.03,

Figure 7: An example line from a weather data file preceded by the many, but thankfully explicit, column headings.

equal distribution of late and not-late ferry trips. This is an incredibly important feature to preserve, since we want to both learn about tardiness and check if the linear classifier produced predicts both scenarios. We stated earlier that it was unknown what amount of data to train on, so the size of the training set here was chosen as a baseline to be examined and experimented with should time permit.

Our data is separated in a **non-uniform** manner. The LibSVM library was used to make the training and testing sets have a roughly proportional, but still random, sampling of the late and not-late trips. This is the default behavior of LibSVM and seemed reasonable at the time. A uniform sampling of the data might be more realistic, but without precedence for it, we feel the distinction between the sampling techniques is only one more in a list of several ways we might tweak the training and analysis to further explore the data.

4 Results

The modeling aspect of the project took place in roughly three phases. As described before, the primary goal models the ferry system and answers the question of whether or not using an SVM to train a linear classifier proves useful for ferry traffic. This leads to the first two phases which involve a first pass and then a serious refining of the methodology. Recall that we consider approximately 80% accurate on the test data to be comparable with Smith et al [SSD08], since like that group, we wish to see if using SVM is a practical method for determining ferry tardiness. A secondary goal is to find which features matter in determining tardiness. While this goal asks questions best answered by a regression analysis, we attempt a crude approximation through empirical methods. The final component is a small experiment to test interesting phenomena arising in the general results. This latter part may be the most interesting of the project, but we begin first with an attempt to use LibSVM.

4.1 The first pass

We first tried determining whether or not a ferry is on time for its departure, leaving arrival for after the process was worked out. A file with all of the labeled training instances was fed into LibSVM, using the default parameters of the kernel. The process took an entire night running on low end hardware, and the results were disappointing to say the least.

After running on the test data, the linear classifier only yielded 60% accuracy. This is significantly beneath the benchmark. The time to run the process was also surprising and led to a more cautious approach for the rest of the work. It became clear at this point how precious each run would be (for time and computing resources).

To improve the performance of the classifier, LibSVM provides an exhaustive guide [CL11] for beginners in the perils of simply “running it”. The directions weren’t followed on the first pass because it was hoped the process would be much simpler just using the default parameters. Although this approach was wrong, it provides anecdotal evidence regarding the importance of being thorough and following the steps in the guide.

Having learned some lessons, the LibSVM guide was followed to the letter on a second pass. The only major change to the process was training on the parameters for the kernel function. As mentioned before LibSVM provides code to perform an exhaustive search for the best fit parameters to use in the function. The subsets for this training come from the training data. Finding the best fit parameters is a long process, sometimes as long as the actual SVM training, so it slows down the process significantly, but as we’re about to discuss the results are well worth it.

4.2 General results

Training the SVM to predict arrival time tardiness proved to be the most successful. In this case, a full trip description (feature vector) included features for the actual departure time, all weather, predicted arrival time, date, vessel name, departure location, and arrival location. (Final results for all runs are displayed in **Figure 8**.) We expect for this to be the problem the SVM can be most accurate about, since really it only needs to understand how long the trip takes a ferry.

The tardiness of full feature instances were able to be predicted with $\approx 87.0\%$ accuracy. For a second pass, this is fantastic to be above the benchmark of 80%. Given that about 87% of trips are on time though, this accuracy shouldn’t be surprising. In fact, we could achieve the same results by guessing all trips are on time, so this model isn’t as exciting as we might like. We’ll examine more ways to check the usefulness of the model in **Section 4.4**.

To tune the linear classifier from our SVM, we should try trimming unnecessary features to find maximum accuracy. Due to time constraints, the trimming we perform here is very coarse. A thorough examination would attempt to examine all combinations of features, but this is beyond the scope of our work, especially in regards to available computing power.

Removing the departure information from the features drops accuracy to 60%, a disappointing result. It may be possible to dramatically improve this result with some of the techniques discussed in **Section 5.1**. To determine whether removing more features made any significant difference, we also removed the weather information from predicting arrivals, yielding the more interesting result which we discuss after departure prediction.

A full departure trip description includes all of the information regarding a trip (vessel name, estimated time of departure, route) and the weather. The prediction results here were not as exciting

Arrival Predictions	
Configuration	Percent accuracy on test data
Full	87%
No Departure	60%
No Weather	88%
No Departure, No Weather	59.4%

Figure 8: Feature configurations and their associated percent accuracies in predicting the timeliness of arrivals convey the empirical importance of the features included.

Departure Prediction Accuracy	
Configuration	Percent accuracy on test data
Full	76.3%
No Weather	75.7%

Figure 9: Feature configurations and their associated percent accuracies in predicting the timeliness of departures convey the empirical importance of the features included.

as with arrivals, but they proved to be more stable. The full departure vectors were predicted with 76.3% accuracy, and removing weather only reduced accuracy to 75.7%. While 76% is not at the benchmark, having experienced the change removing features and adjusting the training sets can have, it seems like more time with the data could yield even better results.

When departure information is still included in arrival predictions, the removal of weather actually improves accuracy resulting in 88%, an improvement of one percentage point. Without the departure information, the percent accuracy decreases by only 0.6 percentage points, a negligible difference at 59.4% accuracy. Weather was included as all but a given for improving the accuracy of the SVM work, but the results just outlined call into question this assumption and warranted further investigation.

4.3 The problem with weather

Since weather made barely any difference in prediction accuracy (even degrading it in some cases), examining possible causes of this unexpected incident became the last focus of the project. By re-examining the assumptions made regarding weather at the start, there are several which we can now dispute:

1. Weather helps improve accuracy.
2. The area is geographically small enough that one weather station should be just fine for all routes.
3. The weather variables dry bulb, wet bulb, dew point, wind speed, wind direction, and station pressure are the right variables to examine.
4. The data retrieved is of a sufficient quality.

Given time constraints and the results of the SVM, it seems assumption 2 is the best to tackle. But what how can we tease apart this assumption? Our strategy was to tackle the effect of the station location by comparing the effect of weather on three routes sufficiently far apart. Running individual SVMs for each route (testing accuracy only on that route's trips). The three routes chosen were Tahlequah-Point Defiance, Kingston-Edmonds, and Coupeville-Port Townsend. Looking at

Prediction accuracies with a *Weather/NoWeather* format

Arrivals	Departures
Tahlequah-Point Defiance	
95.2%/93.9%	76.1%/75.3%
Kingston-Edmonds	
85.2%/94.7%	82.7%/82.8%
Coupeville-Port Townsend	
88.3%/87.7%	68.8%/69.4%

Figure 10: The arrival and departure time prediction accuracies for individual routes when including and excluding weather.

Figure 5, we can see the furthest route, Coupeville-Port Townsend, is nearly 100 miles from the weather station near Point Defiance. These three were deemed to be closest, middlemost, and furthest from the weather station, respectively, but also have enough trips to warrant an individual SVM.

The results of the experiment are mixed and can be viewed in full with **Figure 10**. We used the full feature vectors (described in **Section 4.2**) for predicting both arrivals and departures, only taking out weather where noted for the comparison. We can see that the removal of weather rarely made a difference in any of the measures, except for the Kingston-Edmonds route when looking at arrivals. In that case the removal of weather was a major boon the accuracy. Otherwise, the results here are inconclusive of the actual efficacy of weather in improving predictions, and if anything, only help confuse the situation.

The robustness of the ferry system provides another possible explanation. Recall we that our definition leads to only about 13% of ferry trips being late, a remarkable feat for the WSDOT. It may simply be that there isn't much of a pattern in regards to weather or other predictable phenomena. We may be looking at a system whose only fault is vessel maintenance, a feature requiring very different data than that which we had access to.

The accuracy of the linear classifiers produced by our SVM runs are promising, but so far we have used accuracy as the only measure of effectiveness. We now turn our attention to confusion matrices, precision, and recall. These features allow us to more thoroughly analyze the first linear classifiers we produced.

4.4 Beyond accuracy

As mentioned before, accuracy is only a crude measure of success. It also binds us to comparing our models with a “naive” model which always guesses not-late and is 87% accurate on the full dataset. We should note, however, that our testing data is split such that about 61% are not-late. This means the naive model is only 61% accurate on this testing data, certainly making our models look a bit better. Even so, it's an incredibly basic model we should strive to improve on or at least understand better.

Investigating the precision and recall of our models each other, with more than accuracy, reveal the nuances of prediction [Elk12]. We use confusion matrices to visualize the results and help calculate the precision and recall. A confusion matrix displays the actual classes against the predicted classes to convey the true positive (late), true negative (not-late), false positive, and false negative results:

Arrivals	
	actual class
predicted class	17,574 (<i>tp</i>) 1,974 (<i>fp</i>)
	4,477 (<i>fn</i>) 25,975 (<i>tn</i>)
Arrivals without weather	
	actual class
predicted class	18,122 (<i>tp</i>) 2,115 (<i>fp</i>)
	3,878 (<i>fn</i>) 25,885 (<i>tn</i>)

Figure 11: The confusion matrices for predicting arrival tardiness with the departure information.

	actual class	
predicted class	True Positive (<i>tp</i>)	False Positive(<i>fp</i>)
	False Negative (<i>fn</i>)	True Negative(<i>tn</i>)

Here we are considering positive to mean a trip flagged as late and negative as not-late. The diagonal of the matrix tells us the trips our classifier correctly predicts, and non-zero values off the diagonal indicate incorrect classifications. Thus, we calculate precision by

$$Precision = \frac{tp}{tp + fp}$$

and recall with

$$Recall = \frac{tp}{tp + fn}.$$

Higher precision corresponds to an ability to return *more relevant results than irrelevant results*, and high recall means an algorithm returns *most of the relevant results*. So with high precision, a trip flagged as late is more likely to be late than not-late. High recall means our model catches most of the late trips. These two measures go beyond just accuracy and tell us more about the usefulness of our model.

First note we are using the testing sets to find this data, so the number of trips we are considering is 50,000. Also, a perfect score for either recall or precision is 1. It’s incredibly unlikely, and probably unreasonable, to expect such a score. In fact, the only reasonable scenario for this would be with an entirely deterministic system. We now look at scores for the weather/no-weather models of arrival departure tardiness from **Section 4.2**. This examination hopefully makes clear the differences beyond accuracy.

We start by examining the arrivals with departure information. Recall the classifier is 87% accurate with weather and 88% without, better on testing data than naive model just discussed, but we consider this model to have an “unfair advantage” by knowing information about departures. **Figure 11** displays this model’s confusion matrix. The high accuracy of this model reveals itself in the low numbers off of the diagonals, but the precision and recall are what we hope differentiates this model from naive guessing.

In **Figure 12**, note the very high precision and recall for both the models with and without weather. The high precision indicates this model returns substantially more late trips as late than falsely identifying not-late trips. The high recall also suggests our model returns most of the late results as late. These are good results, but it remains that this model requires the departure information, something a user might not have if they look for a prediction before the departure information becomes available. So, let us turn to the arrivals model without departure information.

Arrivals: precision and recall

Classifier	Precision
Arrivals	0.899
Arrivals without weather	0.895
	Recall
Arrivals	0.797
Arrivals without weather	0.824

Figure 12: Precision and recall for the arrivals model with departure information.

Arrivals without departures

	actual class	
predicted class	6,146 (<i>tp</i>)	4,692 (<i>fp</i>)
	15,905 (<i>fn</i>)	23,257 (<i>tn</i>)

Arrivals without departures and weather

	actual class	
predicted class	8,696 (<i>tp</i>)	6,969 (<i>fp</i>)
	13,355 (<i>fn</i>)	20,980 (<i>tn</i>)

Figure 13: The confusion matrices for predicting arrival tardiness without the departure information.

The arrivals model without departure information achieved 60% accuracy with weather and 59.4% without. These are significantly lower accuracies than our model with departure information, but it's possible the precision and recall values for the model make it competitive in some sense. (Especially since it can be used long before the departure information becomes available to a user.) **Figure 13** contains the confusion matrix for this arrivals model.

The lower accuracy of this model comes through in the high numbers off of the diagonals. Otherwise, there isn't anything immediately noteworthy regarding this model. Neither the true negatives nor true positives stand out as high, so the model probably isn't better at classifying late events or classifying not-late events. What can we learn from the precision and recall though? For this information, we look to **Figure 14**.

This time, the scores are significantly lower, expected given the low accuracy of this model. The precisions are almost at 50%, indicating the model is confusing not-late and late events. Additionally, the recalls are below 50%, meaning a large number of late events are being missed. The low recall is most troubling: users would be expecting many boats to be late when they aren't. In a way, this model isn't aggressive enough. The fact that the model without weather has a significantly higher recall is interesting. Weather continues to yield bizarre effects, possibly suggesting weather information confuses the ability to pick out late events. This is definitely worthy of further investigation, but with scores so low we can label this model defunct and requiring significant further work. We still have the departures model, however, for which the confusion matrix sits in **Figure 15**.

Arrivals without departures: precision and recall

Classifier	Precision
Arrivals without departures	0.588
Arrivals without departures and weather	0.555
	Recall
Arrivals without departures	0.279
Arrivals without departures and weather	0.394

Figure 14: Precision and recall for the arrivals model without departure information.

Departures

	actual class	
predicted class	1,466 (<i>tp</i>)	993 (<i>fp</i>)
	10,853 (<i>fn</i>)	36,688 (<i>tn</i>)

Departures without weather

	actual class	
predicted class	158 (<i>tp</i>)	445 (<i>fp</i>)
	12,161 (<i>fn</i>)	37,236 (<i>tn</i>)

Figure 15: The confusion matrices for predicting departure tardiness.

Departures: precision and recall

Classifier	Precision
Full departure	0.597
Departure without weather	0.262
	Recall
Full departure	0.119
Departure without weather	0.013

Figure 16: Precision and recall for the departure model.

We can see the classifier for departures without weather predicts quite a few more trips as not-late. Remembering that the departure classifier with weather is 76.3% accurate and the classifier without weather is 75.7% accurate, perhaps the weather information is only good for the relatively small number of trips that are actually late. Having investigated the other confusion matrices, this seems like a possible link: weather is important only for the “rare” late cases. It’s an interesting idea for later study. For now, we turn to the precision and recall values of the model in **Figure 16**.

We can see the full departure classifier actually scores much higher in both precision and recall. The recall score means the classifier flags about 12% of late trips, and the approximately 60% precision means the classifier flags mostly late trips as late, rather than indicating many not-late trips as late. The recall suggests we may be missing some indicator of what makes a trip late. In the case of precision, the number is encouraging, indicating we aren’t confusing late and not-late trips, for the most part. The numbers are so much lower than for arrivals though, even our low accuracy model without departure information, this departure model isn’t looking quite as good. Such a low recall means a user of this model might feel like the application is always guessing not-late. Obviously, this model needs further tuning.

To tune these parameters, we could try any number of approaches (many outlined in **Section 5.1**). The arrival models are much more interesting than their accuracy indicates, especially given how low the departures model scores. After tweaking and testing more, we might find the arrivals are generally more predictable than departures or there just isn’t a way to consistently classify those fringe tardy trips. Regardless, the first pass at these models is encouraging.

5 Closing thoughts

We met the original goal of the project, to develop a ferry system model using AI techniques, and further investigated the behavior of the system. We beat the initial benchmark of 80% accuracy for a linear classifier produced by an SVM, and we concluded the performance could probably be improved on given more time and experience tuning the process. The exploration of weather’s relationship to

the ferry system also proved to be more complex than expected, clouding the assumption of there being a significant link between the two. We also examined the precision and recall, discovering mixed results between the models. Our most effective model, which had the advantage of departure information, exhibited high precision and recall, indicating it may be worth further study as an improvement upon the naive guess of always not-late. The remaining questions and further work are now even greater in scope than when the project began.

5.1 Improvements to the data and model analysis

Most of the percent accuracies we analysed were short of the 80% benchmark, but some were very close. More time to tune the features used in prediction, check the data quality, and experiment with different training set sizes would likely yield even better predictions. The latter goal is of particular interest. Since no definitive guide to using SVM was found before finishing the project, the size of training sets was kept the same throughout but may be a poorly informed choice. A common criticism of SVMs, and AI in general, is that training on too large of a dataset can lead to overfitting without realising. This means the machine only accurately predicts events similar to the training data and can't generalize beyond. With a chance to train on smaller sets of data, we might find overfitting hampers the accuracy we found with larger training sets.

Recall the variance found among training an SVM over individual routes in the weather investigation. We discovered that individual routes varied widely, suggesting a “master” SVM for all routes may not be appropriate. We would like to assume that an SVM could tease apart the relationships between routes and tardiness, but there may be too many conflicts between routes in terms of which features matter. For instance, a short route between Clinton and Mukilteo covers mostly open water and might not be very susceptible to weather in comparison to the island crowded route from Anacortes to Lopez. A thorough comparison of how individual SVMs perform for the routes might make this approach even more appealing or lend insight into a better choice.

Section 4.4 explored confusion matrices as a means of slicing and dicing the finer details of our models' performance. That discussion only focused on the models we produced and not necessarily how to improve them. Using information from the confusion matrix, we can explore the receiver operating characteristic (more simply called a ROC curve). This graphical plot depicts the true positive rate against the false positive rate to help in selecting a “best” model. To use a ROC curve, we plot the true positive and false positive rates as parameters are tweaked to produce different models.

Analysing a ROC curve gives us a means of understanding the trade-offs between precision and recall. Many cases require a trade-off in one or the other, a trait we haven't examined in this project. (As an example, consider a surgeon deciding which tissue to remove when dealing with cancer. Removing too little may increase precision, more cells removed are cancerous, but significantly decrease recall, getting too few cancerous cells. The other way around also has evident advantages and disadvantages.) If we did examine the trade-offs in the ferry system, it could be our model simply can't win in both metrics. This information would actually tell us a great deal about the usefulness of any model for this situation!

The end product of our SVM trainings worked in the ways we hoped: seeming to be useful for predicting ferry timeliness and as an investigative tool for the WSDOT ferry system. Still, the project is far from complete. Just changing the training/testing sampling to be uniform, for instance, could

greatly influence the success of our classifier, but it would also be interesting just to compare the results of using a different sampling technique.

5.2 Extensions

The SVM training produced useful linear classifiers, but a framework for average users could still be built. Whether this would be a web or mobile application is unfortunately restricted by the use of LibSVM. Since the only implementation is in C, a web application querying some server seems like the most straightforward approach. There are other possibilities since LibSVM has bindings in many languages, but the production of a real user interface and seeing it used would certainly validate the usefulness of predicting ferry times for the riders.

The SVM training itself could also be restructured. Several times, the lack of production time inhibited the ability to fully explore the data, but beyond this we could improve the definition of tardy. A multiclass SVM allows us to have degrees of tardiness (e.g. three minutes, five minutes, or ten minutes late). LibSVM supports this sort of SVM, a fairly natural extension of the two class case. Moreover, LibSVM could be used for a regression analysis of the features. This might be the best way to determine which features matter in regards to tardiness.

5.3 Conclusion

The use of an SVM for predicting ferry times novelly makes use of the large data sets available. This tool from artificial intelligence enables us to predict, with over 75% accuracy in most cases, the tardiness of ferries in the complex system, and it brings us insight into how important various features are when predicting this tardiness. While we may not fully understand the strange behavior exhibited by including weather in the model, the base assumption of weather's importance no longer appears valid. Further tuning of the SVM process may yield even greater insight and perhaps better prediction results. In general, however, the results of the project are promising and suggest support vector machines are a useful approach to modeling the ferry system.

References

- [ABR64] A Aizerman, Emmanuel M Braverman, and LI Rozoner, *Theoretical foundations of the potential function method in pattern recognition learning*, Automation and remote control **25** (1964), 821–837.
- [BFSS03] Evgeny Byvatov, Uli Fechner, Jens Sadowski, and Gisbert Schneider, *Comparison of support vector machine and artificial neural network systems for drug/nondrug classification*, Journal of Chemical Information and Computer Sciences **43** (2003), no. 6, 1882–1889.
- [CL11] Chih-Chung Chang and Chih-Jen Lin, *Libsvm: a library for support vector machines*, ACM Transactions on Intelligent Systems and Technology (TIST) **2** (2011), no. 3, 27.
- [Elk12] Charles Elkan, *Evaluating Classifiers*, <http://cseweb.ucsd.edu/users/elkan/250Bwinter2012/classifiereval.pdf>, January 2012.
- [OA12] National Oceanic and Atmospheric Administration, *NCDC: Quality Controlled Local Climatological Data*, <http://cdo.ncdc.noaa.gov/qclcd/QCLCD>, October 2012.
- [oT12a] Washington State Department of Transportation, *Washington State Ferries Our Fleet*, <http://www.wsdot.wa.gov/Ferries/yourwsf/ourfleet/>, October 2012.
- [oT12b] ———, *WSDOT - Ferries - VesselWatch*, <http://www.wsdot.com/ferries/vesselwatch/Default.aspx>, January 2012.
- [SSD08] David A Smith, Lance Sherry, and G Donohue, *Decision support tool for predicting aircraft arrival rates, ground delay programs, and airport delays from weather forecasts*, International Conference on Research in Air Transportation (ICRAT-2008), 2008.
- [Was12] Washington State Department of Transportation, *WSF Traffic Statistics*, http://www.wsdot.wa.gov/ferries/traffic_stats/, October 2012.